

# Artificial Intelligence

## Machine Learning

## Neural Networks




# Neural Networks

---

- Algorithms that try to mimic how the brain functions.
- Worked extremely well to recognize:
  1. handwritten characters (LeCun et al. 1989),
  2. spoken words (Lang et al. 1990),
  3. faces (Cottrel 1990)
- Extensively studied in the 1990's with a moderate success.
- Now back with lots of success with deep learning thanks to the algorithmic and computational progress.
- The first algorithm used was the Perceptron (Resenblatt 1959).

# Perceptron

---



Given  $n$  examples and  $d$  features.

$$f(x_i) = \text{sign}\left(\sum_{j=0}^d w_j x_{ij}\right)$$

# Perceptron expressiveness

---

- Consider the perceptron with the step function.
- Idea: Iterative method that starts with a random hyperplane and adjust it using your training data.
- It can represent Boolean functions such as AND, OR, NOT but not the XOR function.
- It produces a linear separator in the input space.

# From perceptron to MLP

---


- The perceptron works perfectly if data is linearly separable. If not, it will not converge.
- Neural networks use the ability of the perceptrons to represent elementary functions and combine them in a network of layers of elementary questions.
- However, a cascade of linear functions is still linear,
- and we want networks that represent highly non-linear functions.

# From perceptron to MLP

---


- Also, perceptron used a **threshold function**, which is undifferentiable and not suitable for gradient descent in case data is not linearly separable.
- We want a function whose output is a linear function of the inputs.
- One possibility is to use the sigmoid function:

$$g(z) = \frac{e^z}{1 + e^z} = \frac{1}{1 + e^{-z}}$$



$$g(z) \rightarrow 1 \text{ when } z \rightarrow +\infty \qquad g(z) \rightarrow 0 \text{ when } z \rightarrow -\infty$$

# Perceptron with Sigmoid



Given  $n$  examples and  $d$  features.

For an example  $x_i$  (the  $i^{th}$  line in the matrix of examples)

$$f(x_i) = \frac{1}{1 + e^{-\sum_{j=0}^d w_j x_{ij}}}$$

# The XOR example

---

Let's try to create a MLP for the XOR function using elementary perceptrons.



# The XOR example

---

Let's try to create a NN for the XOR function using elementary perceptrons.

First observe:

$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g(10) = 0.99995$$

$$g(-10) = 0.00004$$

Let's consider that: For  $z \geq 10$ ,  $g(z) \rightarrow 1$ . For  $z \leq -10$ ,  $g(z) \rightarrow 0$ .

# The XOR example


---

First what is the perceptron of the OR?

# The XOR example

---

$x_1$	$x_2$	$x_1$ <b>OR</b> $x_2$	$g(z)$
0	0	0	$g(w_0 + w_1x_1 + w_2x_2) = g(-10)$
0	1	1	$g(10)$
1	0	1	$g(10)$
1	1	1	$g(10)$



# The XOR example



---

Similarly, we obtain the perceptrons for the AND and NAND:

# The XOR example

---

Similarly, we obtain the perceptrons for the AND and NAND:



Note: how the weights in the NAND are the inverse weights of the AND.

# The XOR example

---

Let's try to create a NN for the XOR function using elementary perceptrons.

$x_1$	$x_2$	$x_1$ <b>XOR</b> $x_2$	$(x_1$ <b>OR</b> $x_2)$ <b>AND</b> $(x_1$ <b>NAND</b> $x_2)$
0	0	0	0
0	1	1	1
1	0	1	1
1	1	0	0

# The XOR example


---

Let's put them together...

# The XOR example

---

Let's put them together...



XOR as a combination of 3 basic perceptrons.



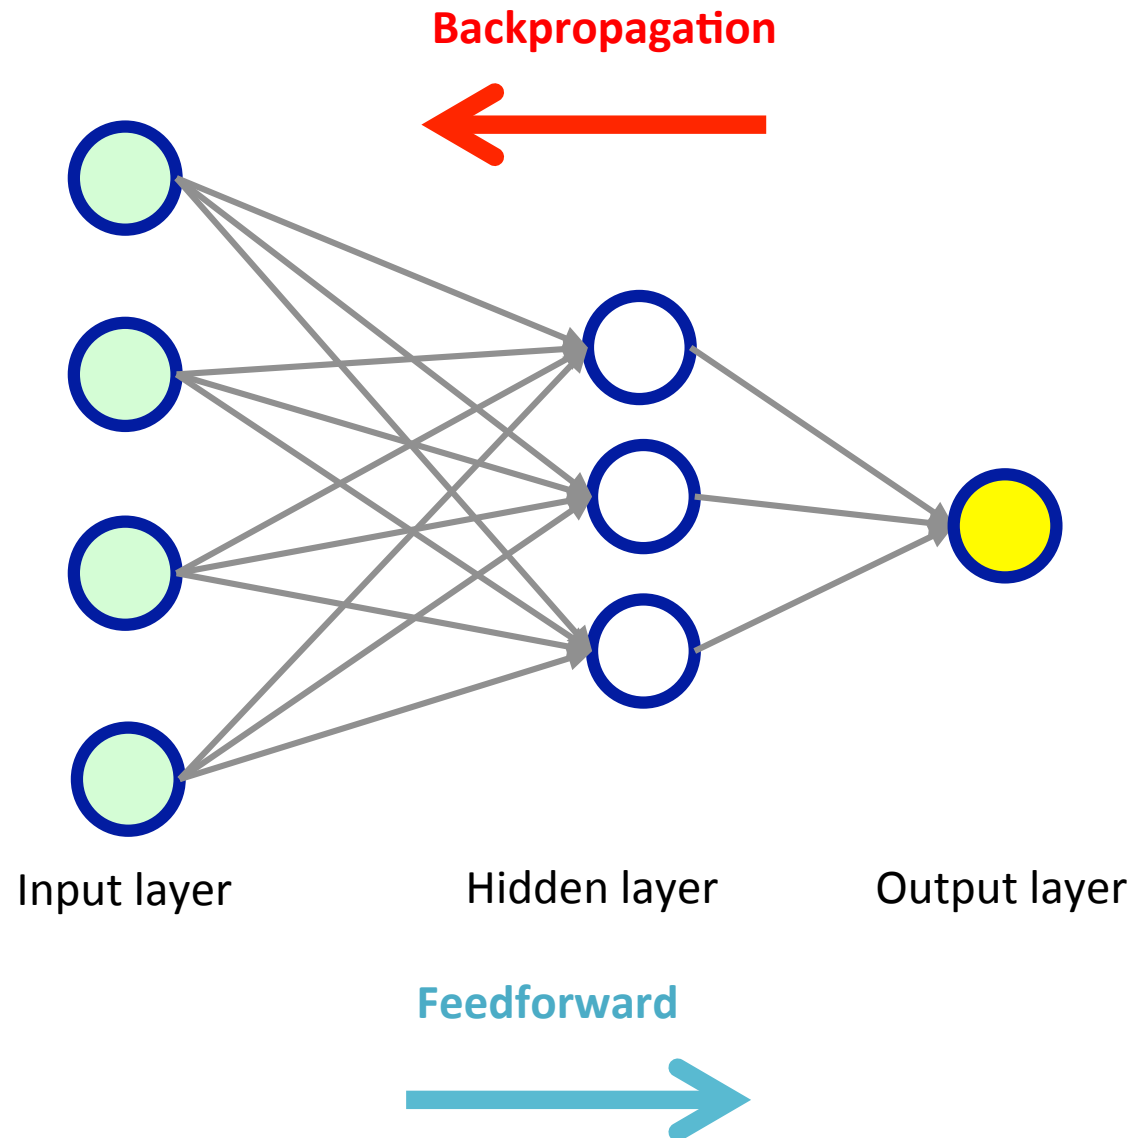
# Backpropagation algorithm

---

- Note: Feedforward NN (as opposed to recurrent networks) have no connections that loop.
- Learn the weights for a multilayer network.
- Backpropagation stands for “backward propagation of errors”.
- Given a network with a fixed architecture (neurons and interconnections).
- Use Gradient descent to minimize the squared error between the network output value  $o$  and the ground truth  $y$ .
- We suppose multiple output  $k$ .
- Challenge: Search in all possible weight values for all neurons in the network.

# Feedforward-Backpropagation

---



# Backpropagation rules

---

- We consider  $k$  outputs
- For an example  $e$  defined by  $(x, y)$ , the error on training example  $e$ , summed over all output neurons in the network is:

$$E_e(w) = \frac{1}{2} \sum_k (y_k - o_k)^2$$

- Remember, gradient descent iterates through all the training examples one at a time, descending the gradient of the error w.r.t. this example.

$$\Delta w_{ij} = -\alpha \frac{\partial E_e(w)}{\partial w_{ij}}$$

# Backpropagation rules

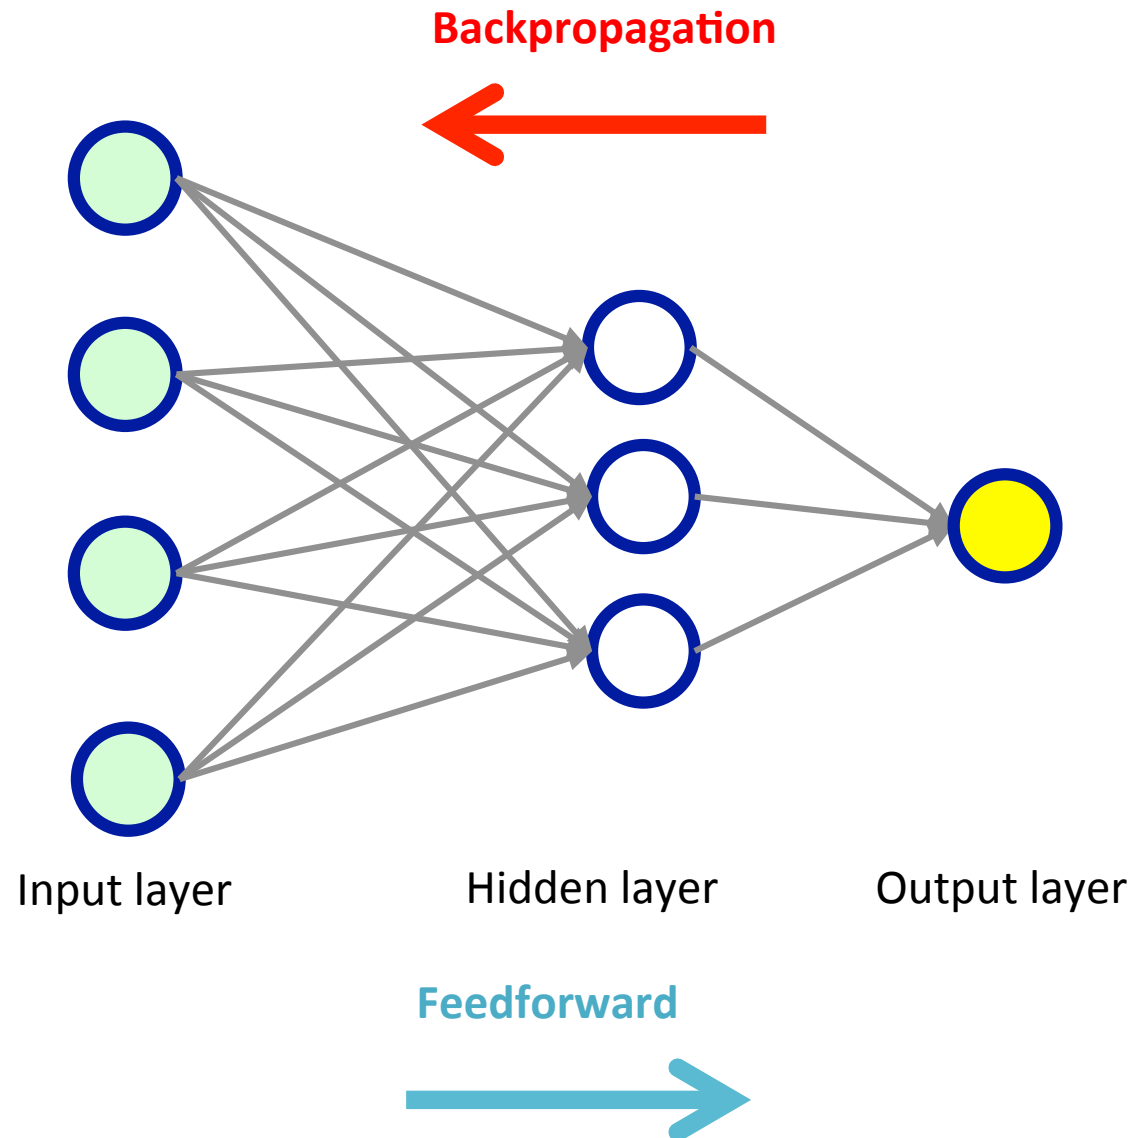
---

## Notations:

- $x_{ij}$ : the  $i^{th}$  input to neuron  $j$ .
- $w_{ij}$ : the weight associated with the  $i^{th}$  input to neuron  $j$ .
- $z_j = \sum w_{ij}x_j$ , weighted sum of inputs for neuron  $j$ .
- $o_j$ : output computed by neuron  $j$ .
- $g$  is the sigmoid function.
- *outputs*: the set of neurons in the output layer.
- $Succ(j)$ : the set of neurons whose immediate inputs include the output of neuron  $j$ .

# Backpropagation notations

---



# Backpropagation rules

---

$$\frac{\partial E_e}{\partial w_{ij}} = \frac{\partial E_e}{\partial z_j} \frac{\partial z_j}{\partial w_{ij}} = \frac{\partial E_e}{\partial z_j} x_{ij}$$

$$\Delta w_{ij} = -\alpha \frac{\partial E_e}{\partial z_j} x_{ij}$$

We consider two cases in calculating  $\frac{\partial E_e}{\partial z_j}$  (let's abandon the index  $e$ ):

- **Case 1:** Neuron  $j$  is an output neuron
- **Case 2:** Neuron  $j$  is a hidden neuron

# Backpropagation rules

---

- Case 1: Neuron  $j$  is an output neuron

$$\frac{\partial E}{\partial z_j} = \frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial z_j}$$

# Backpropagation rules

---

- Case 1: Neuron  $j$  is an output neuron

$$\frac{\partial E}{\partial z_j} = \frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial z_j}$$

$$\frac{\partial E}{\partial o_j} = \frac{\partial}{\partial o_j} \frac{1}{2} \sum_k (y_k - o_k)^2$$

$$\frac{\partial E}{\partial o_j} = \frac{\partial}{\partial o_j} \frac{1}{2} (y_j - o_j)^2$$

$$\frac{\partial E}{\partial o_j} = \frac{1}{2} 2 (y_j - o_j) \frac{\partial (y_j - o_j)}{\partial o_j}$$

$$\frac{\partial E}{\partial o_j} = -(y_j - o_j)$$



# Backpropagation rules

---

- Case 1: Neuron  $j$  is an output neuron

$$\frac{\partial E}{\partial z_j} = \frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial z_j}$$

$$\frac{\partial E}{\partial o_j} = \frac{\partial}{\partial o_j} \frac{1}{2} \sum_k (y_k - o_k)^2$$

$$\frac{\partial E}{\partial o_j} = \frac{\partial}{\partial o_j} \frac{1}{2} (y_j - o_j)^2$$

$$\frac{\partial E}{\partial o_j} = \frac{1}{2} 2 (y_j - o_j) \frac{\partial (y_j - o_j)}{\partial o_j}$$

$$\frac{\partial E}{\partial o_j} = -(y_j - o_j)$$

We have:  $o_j = g(z_j)$

$$\frac{\partial o_j}{\partial z_j} = \frac{\partial g(z_j)}{\partial z_j}$$

$$\frac{\partial o_j}{\partial z_j} = o_j(1 - o_j)$$

# Backpropagation rules

---

$$\frac{\partial E}{\partial z_j} = -(y_j - o_j)o_j(1 - o_j)$$

$$\Delta w_{ij} = \alpha(y_j - o_j)o_j(1 - o_j)x_{ij}$$

We will note

$$\delta_j = -\frac{\partial E}{\partial z_j}$$

$$\Delta w_{ij} = \alpha \delta_j x_{ij}$$

# Backpropagation rules

---

- Case 2: Neuron  $j$  is a hidden neuron

$$\frac{\partial E}{\partial z_j} = \sum_{k \in \text{succ}\{j\}} \frac{\partial E}{\partial z_k} \frac{\partial z_k}{\partial z_j} = \sum_{k \in \text{succ}\{j\}} -\delta_k \frac{\partial z_k}{\partial z_j}$$

$$\frac{\partial E}{\partial z_j} = \sum_{k \in \text{succ}\{j\}} -\delta_k \frac{\partial z_k}{\partial o_j} \frac{\partial o_j}{\partial z_j}$$

$$\frac{\partial E}{\partial z_j} = \sum_{k \in \text{succ}\{j\}} -\delta_k w_{jk} \frac{\partial o_j}{\partial z_j}$$

$$\frac{\partial E}{\partial z_j} = \sum_{k \in \text{succ}\{j\}} -\delta_k w_{jk} o_j (1 - o_j)$$

$$\delta_j = -\frac{\partial E}{\partial z_j} = o_j (1 - o_j) \sum_{k \in \text{succ}\{j\}} \delta_k w_{jk}$$

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

1. Create\_feedforward\_network ( $n_i$ ,  $n_h$ ,  $n_o$ )

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

1. Create\_feedforward\_network ( $n_i$ ,  $n_h$ ,  $n_o$ )
2. Initialize all weights to a small random number (e.g., in  $[-0.2, 0.2]$ )

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

1. Create\_feedforward\_network ( $n_i$ ,  $n_h$ ,  $n_o$ )
2. Initialize all weights to a small random number (e.g., in  $[-0.2, 0.2]$ )
3. Repeat until convergence
  - (a) For each training example  $(x, y)$



# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

1. Create\_feedforward\_network ( $n_i$ ,  $n_h$ ,  $n_o$ )
2. Initialize all weights to a small random number (e.g., in  $[-0.2, 0.2]$ )
3. Repeat until convergence
  - (a) For each training example  $(x, y)$ 
    - i. **Feed forward:** Propagate example  $x$  through the network and compute the output  $o_j$  from every neuron neuron.

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

1. Create\_feedforward\_network ( $n_i$ ,  $n_h$ ,  $n_o$ )
2. Initialize all weights to a small random number (e.g., in  $[-0.2, 0.2]$ )
3. Repeat until convergence
  - (a) For each training example  $(x, y)$ 
    - i. **Feed forward:** Propagate example  $x$  through the network and compute the output  $o_j$  from every neuron neuron.
    - ii. **Propagate backward:** Propagate the errors backward.

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

1. Create\_feedforward\_network ( $n_i$ ,  $n_h$ ,  $n_o$ )
2. Initialize all weights to a small random number (e.g., in  $[-0.2, 0.2]$ )
3. Repeat until convergence
  - (a) For each training example  $(x, y)$ 
    - i. **Feed forward:** Propagate example  $x$  through the network and compute the output  $o_j$  from every neuron neuron.
    - ii. **Propagate backward:** Propagate the errors backward.
      - Case 1** For each output neuron  $k$ , calculate its error
$$\delta_k = o_k(1 - o_k)(y_k - o_k)$$

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .

**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

1. Create\_feedforward\_network ( $n_i$ ,  $n_h$ ,  $n_o$ )
2. Initialize all weights to a small random number (e.g., in  $[-0.2, 0.2]$ )
3. Repeat until convergence
  - (a) For each training example  $(x, y)$ 
    - i. **Feed forward:** Propagate example  $x$  through the network and compute the output  $o_j$  from every neuron neuron.
    - ii. **Propagate backward:** Propagate the errors backward.
      - Case 1** For each output neuron  $k$ , calculate its error
$$\delta_k = o_k(1 - o_k)(y_k - o_k)$$
      - Case 2** For each hidden neuron  $h$ , calculate its error
$$\delta_h = o_h(1 - o_h) \sum_{k \in Succ(h)} w_{hk} \delta_k$$

# Backpropagation algorithm

---

**Input:** training examples  $(x, y)$ , learning rate  $\alpha$  (e.g.,  $\alpha = 0.1$ ),  $n_i$ ,  $n_h$  and  $n_o$ .


**Output:** a neural network with one input layer, one hidden layer and one output layer with  $n_i$ ,  $n_h$  and  $n_o$  number of neurons respectively and all its weights.

1. Create\_feedforward\_network ( $n_i$ ,  $n_h$ ,  $n_o$ )
2. Initialize all weights to a small random number (e.g., in  $[-0.2, 0.2]$ )
3. Repeat until convergence
  - (a) For each training example  $(x, y)$ 
    - i. **Feed forward:** Propagate example  $x$  through the network and compute the output  $o_j$  from every neuron neuron.
    - ii. **Propagate backward:** Propagate the errors backward.
      - Case 1** For each output neuron  $k$ , calculate its error
$$\delta_k = o_k(1 - o_k)(y_k - o_k)$$
      - Case 2** For each hidden neuron  $h$ , calculate its error
$$\delta_h = o_h(1 - o_h) \sum_{k \in Succ(h)} w_{hk} \delta_k$$
    - iii. Update each weight  $w_{ij} \leftarrow w_{ij} + \alpha \delta_j x_{ij}$

# Observations

---

- Convergence: small changes in the weights
- There are other activation functions. Hyperbolic tangent function, is practically better for NN as its outputs range from -1 to 1.




$$g(x) = \text{sigmoid}(x) = \frac{e^{kx}}{1+e^{kx}} \text{ for } k = 1, k = 2, \text{ etc.}$$

$$g(x) = \text{tanh}(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \text{ (It is a rescaling of the logistic sigmoid function!).}$$


# Multi class case etc.

---



# Multi class case etc.

---



- Nowadays, networks with more than two layers, *a.k.a.* deep networks, have proven to be very effective in many domains.
- Examples of deep networks: restricted Boltzman machines, convolutional NN, auto encoders, etc.




# MNIST database

---

<http://yann.lecun.com/exdb/mnist/>

- The MNIST database of handwritten digits
- Training set of 60,000 examples, test set of 10,000 examples
- Vectors in  $\mathbb{R}^{784}$  (28x28 images)
- Labels are the digits they represent
- Various methods have been tested with this training set and test set



- Linear models: 7% - 12% error
- KNN: 0.5%- 5% error
- Neural networks: 0.35% - 4.7% error
- Convolutional NN: 0.23% - 1.7% error

# Demo: Tensorflow

---

<http://playground.tensorflow.org/>

- Open source software to play with neural networks in your browser.
- The dots are colored orange or blue for positive and negative examples.
- It's possible to choose the activation function, architecture, rate etc.
- Very well done! Let's check it out!

# Credit

---

- Machine Learning 1997. T. Mitchell.
- Andrew Ng's lecture notes.
- <http://playground.tensorflow.org/>